

GWD-R-P
NML-WG
nml-wg@ogf.org

Jeroen van der Ham, University of Amsterdam
Freek Dijkstra, SARA
Jason Zurawski, Internet2
Martin Swany, Indiana University

February 2012

Network Markup Language Base Schema version 1

Status of This Document

Community Practice (CP)

Copyright Notice

Copyright © Open Grid Forum (2008-2012). Some Rights Reserved. Distribution is unlimited.

Abstract

This document describes a normative schema which allows the description of a computer network topology.

Contents

Abstract	1
Contents	1
1 introduction	3
1.1 Notational Conventions	3
2 NML Topology Schema	4
2.1 Network	4
2.2 Network Object	4
2.3 Node	6
2.4 Port	6
2.5 Link	7
2.6 Service	7
2.7 Group	8
2.7.1 Bidirectional Link	8

2.7.2	Bidirectional Port	8
2.7.3	Topology	8
2.7.4	Domain	9
2.7.5	Network	9
2.8	Adaptations	9
2.9	Relation	9
3	Identifiers	11
3.1	Object Identifiers	11
3.2	Instance Identifiers	11
3.2.1	Lexical Equivalence	11
3.2.2	Further Restrictions	12
3.2.3	Interpreting Identifiers	12
3.2.4	Network Object Attribute Change	12
4	Examples	13
5	Security Considerations	14
6	Glossary	15
7	Contributors	16
8	Acknowledgments	16
9	Intellectual Property Statement	17
10	Disclaimer	17
11	Full Copyright Notice	17
References	18
Normative References	18
Informative References	18

1 introduction

This document describes the object model of the Network Markup Language. These basic objects may be extended, or sub-classed, to represent technology specific classes. These basic objects and extended objects will also be representable in multiple syntaxes, including at least XML and RDF.

1.1 Notational Conventions

The keywords “MUST”, “MUST NOT”, “REQUIRED”, “SHALL”, “SHALL NOT”, “SHOULD”, “SHOULD NOT”, “RECOMMENDED”, “MAY”, and “OPTIONAL” are to be interpreted as described in [RFC 2119].

2 NML Topology Schema

The NML Topology schema describes an information model with elements and their relations that describe computer networks. This schema is kept intentionally general, with provisions to extend the base schema to describe layer-specific information.

The URI of the class-objects will become `http://http://schemas.ogf.org/nml/base/yyyy/mm/` where `yyyy/mm` is the four digit year and two digit month of the publication of the schema document.

2.1 Network

This document explicitly does not try to provide a definition of the term ‘Network’. The working group has discussed the possible meanings of this term and in the end we were forced to conclude that is not possible to provide a workable precise definition for the term *Network*. The term *Network* has become so widely used for so many diverse meanings that it is impossible to create a definition that everyone can agree on, while still expressing something useful.

Figure 1 shows an overview of all the objects in the NML schema in a UML class diagram. The figure also shows the relations between the objects, and their cardinalities. In the sections below we discuss each of the elements of the schema.

2.2 Network Object

The basic abstract element of the schema is the *Network Object*. Other basic elements inherit from it. The *Network Object* can have a *Location*, can be related to other instances via *Relations* and can have a *Lifetime*. Every Network Object MUST have an *id* attribute, which MUST be a unique URI. These characteristics are inherited by the subclasses of the *Network Object* class.

The base *Network Object* has three related objects that describe the *Network Object* and its relationships:

- Location
- Lifetime
- Relation

The location of an object in the physical world can be described using the *Location* object. The actual location is then described using properties of the *Location* object. The Location and a Network object are related to each other using the *locatedAt* relationship.




Figure 1: A UML class diagram of the objects in the NML schema and their relations

All *Network Objects* can potentially have a *Lifetime*, that consists of vector of *time* elements, which contain a start time and an end time.

The Relations between different network objects are represented using relation objects. These are discussed in more detail in section 2.9.

The base *Network Object* is subclassed into the top-level topology components, that are sufficient to cover the description of networks. The top-level network objects in this schema are:

- Node
- Port
- Link
- Service
- Group

These objects are described in more detail below.

2.3 Node

A *Node* is generally a device connected to, or part of, the network. A *Node* does not necessarily correspond to a physical machine. It MAY be a virtual device or a group of devices. In this case, the *implementedBy* relation can be used to describe the virtualization.

A *Node* is connected to the network by its *Ports*. A *Node* can provide *Services*.

The Relations of *Node*:

- A *Node* MAY share a *hasPort* relation with one or more *Ports*.
- A *Node* MAY share a *locatedAt* relation with one *Location*
- A *Node* MAY share an *implementedBy* relation with one or more *Nodes*.

2.4 Port

A *Port*, or interface, connects a *Node* or *Group* to the rest of the network. A *Port* object is unidirectional.

A *Port* is related to zero or one *Node* or *Group*, and also has a relation with zero, or one (uni-directional) *Links*.

To adapt traffic to and from different layers, an *adaptation* relation is used. An adaptation consists of two unidirectional components, an *adaptationSink* to go from a server layer port to a client layer port, and an *adaptationSource* for the other direction.

A Port can have up to two adaptation sink relations, one for a server layer, and one for a client layer. The same applies for adaptation source relations.

Relations of Port:

- A *Port* MAY have one *adaptationSink* relations.
- A *Port* MAY have one *adaptationSource* relations.
- A *Port* MAY have one *source* or *sink* relation with a *Unidirectional Link*.

2.5 Link

A *Link* object describes that there is a unidirectional connection from one *Port* to another. These ports are identified using the *source* and *sink* relationships.

A *Link* MUST have an attribute *type* which is either *Link* or *Crossconnect*. When the type is *Crossconnect*, the source and sink MUST be part of the same node.

A *Link* SHOULD have a *capacity* attribute which describes the capacity of the link in bytes per second. This value SHOULD correspond to the actual amount of data that can be transported over the link, excluding overhead.

Relations of Link:

- A *Link* MAY have a *source* relation with one *Port*.
- A *Link* MAY have a *sink* relation with one *Port*.
- A *Link* SHOULD have a *capacity* attribute which describes the capacity of the link in *bytes per second*.

2.6 Service

A *Service* object describes a certain capability being offered by a Network Object. The key idea that this is a generic container representing any service that a network-centric user or agent might want to discover and use. Below we describe some example Services, but others are also possible.

SwitchingMatrix describes the ability of a network object to create cross connects between its different ports.

Configured cross-connects SHOULD be described using the *switchedTo* relation.

Adaptation describes that ports on different layers within a node can possibly be connected together to form a connection or cross connect on a different layer. The Adaptation service MUST define both its client and server layer.

Once this is implemented it is described using the *adaptationSource* and *adaptationSink* relations between ports.

Measurement Point services are an essential component of network measurement services like perfSONAR.

2.7 Group

To describe collections of network objects, there is a group element. Any element defined above can be part of a group, including another group.

We also define a set of special groups:

- Bidirectional Link
- Bidirectional Port
- Topology
- Domain
- Network

2.7.1 Bidirectional Link

A *Bidirectional Link* is a special group of two (unidirectional) *Links* together forming a bidirectional link between two ports.

2.7.2 Bidirectional Port

A *Bidirectional Port* is a special group of two (unidirectional) *Ports* together forming a bidirectional representation of a physical or virtual port.

2.7.3 Topology

A *Topology* is a set of Network Objects and the links connecting them.

2.7.4 Domain

A *Domain* represents an administrative domain as a collection of resources part of that Domain.

2.7.5 Network

A *Network* is an unordered collection of Network Objects managed under the same shared mechanism.

2.8 Adaptations

Adaptations are defined in three different ways:

AdaptationType for the generic type of Adaptations

AdaptationService describes a capability of performing an Adaptation (instance of a Service)

Adaptation instance of an abstract AdaptationType to describe the adaptation being performed in a node.

Adaptation is ‘Actual data transport function where data of one port is embedded in the data of another port.’

2.9 Relation

Relations describe how different network objects can be combined to form a network topology description. The relations have been described above, but for ease of reference we also give a full list and definition here (in alphabetical order):

adaptation is the bi-directional equivalent of the source and sink adaptation combination.

adaptationSink goes from a server layer port to a client layer port, describing the way that data is passed between these two layers.

adaptationSource is the reverse of the sink adaptation, i.e. it goes from a client layer port to a server layer port.

atLayer is a relation between a port and a layer to describe the layer at which the port operates.

hasPort describes the relation between a node or group and a port.

hasService describes the relation between a node or group and a service that it provides.

implementedBy is a relation from a node to a node, describing that the source node is a virtualized node on the sink node of the relation.

locatedAt is a relation between a network object and a location object.

sink is the connection between the end of the link and the sink port.

source describes the connection of a port to the source port of the link.

switchedTo defines a relation between two ports, meaning that traffic from the source port is automatically forwarded to the destination port.

3 Identifiers

3.1 Object Identifiers

The namespace for the class objects defined in this document is `http://schemas.ogf.org/nml/base/2013/`
TODO: change to correct year and month of the schema.

All objects and attributes defined in this document reside in this namespace. For example, the link object is identified by `http://schemas.ogf.org/nml/2013/10/base/link`

3.2 Instance Identifiers

Section 2.2 requires that instances of Network Objects **MUST** have an *id* attribute, which **MUST** be a unique URI.

Implementations that receive a network topology description **MUST** be prepared to accept any valid URI as an identifier.

Implementations that publish a network topology description instance identifiers **MAY** adhere to the syntax of Global Network Identifiers as defined in [URN-OGF-NETWORK], which ensures global uniqueness and that easy recognition of Network Object instances.

Two different Network Objects instance **MUST** have two different identifiers.

Once an identifier is assigned to a resource, it **MUST NOT** be re-assigned to another resource.

A URI **MAY** be interpreted as an International Resource Identifier (IRI) for display purposes, but URIs from external source domains **MUST NOT** be IRI-normalised before transmitting to others.

3.2.1 Lexical Equivalence

Two identifier are lexical equivalent if they are binary equivalent after case-normalisation.

No interpretation of percent-encoding or PUNYCODE decoding should take place.

For the purpose of equivalence comparison, any possible fragment part or query part of the URI is considered part of the URI.

For example the following identifiers are equivalent:

- 1 - `urn:ogf:network:example.net:2012:local_string_1234`
- 2 - `URN:OGF:network:EXAMPLE.NET:2012:Local_String_1234`

while the following identifiers are not equivalent (in this case, the percentage encoding even make URI #3 an invalid Global Network Identifier.):

- 1 - urn:ogf:network:example.net:2012:local_string_1234
- 3 - urn:ogf:network:example.net:2012:local%5Fstring%5F1234

3.2.2 Further Restrictions

An assigning organisation **MUST NOT** assign Network Object Identifier longer than 255 characters in length.

Parsers **MUST** be prepared to accept identifiers of up to 255 characters in length.

A Parser **SHOULD** verify if an identifier adheres to the general URI syntax rules, as specified in RFC 3986 [RFC 3986].

Parsers **SHOULD** reject identifiers which do not adhere to the specified rules. A parser encountering an invalid identifier **SHOULD** reply with an error code that includes the malformed identifier, but **MAY** accept the rest of the message, after purging all references to the Network Object with the malformed identifier.

3.2.3 Interpreting Identifiers

A Network Object identifier **MUST** be treated as a opaque string, only used to uniquely identify a Network Object. The local-part of a Global Network Identifier **MAY** have certain meaning to it's assigning organisation, but **MUST NOT** be interpreted by any other organisation.

3.2.4 Network Object Attribute Change

A Network Object may change during its lifetime. If these changes are so drastic that the assigning organisation considers it a completely new Network Object, the assigning organisation should be assigned a new identifier. In this case, other organisations **MUST** treat this object as completely new Network Resource.

If the assigning organisation considers the changes are small, it **MUST** retain the same identifier for the Network Object, and use some mechanism to signal it's peers of the changes in the attributes of the Network Object.

4 Examples

5 Security Considerations

There are important security concerns associated with the generation and distribution of network topology information. For example, ISPs frequently consider network topologies to be proprietary. We do not address these concerns in this document, but implementers are encouraged to consider the security implications of generating and distributing network topology information.

6 Glossary

7 Contributors

Jeroen J. van der Ham (Editor)

Faculty of Science, Informatics Institute, University of Amsterdam
Science Park 904, 1098 XH Amsterdam
The Netherlands
Email: vdham@uva.nl

Freek Dijkstra

SARA
Science Park 140, 1098 XG Amsterdam
The Netherlands
Email: freek.dijkstra@sara.nl

Jason Zurawski

Internet2
1150 18th Street, NW
Suite 1020
Washington, DC 20036
USA
Email: zurawski@internet2.edu

D. Martin Swany

Indiana University, School of Informatics and Computing
Lindley Hall Room 215
150 S. Woodlawn Avenue
Bloomington, Indiana 47405-7104
USA
Email: swany@iu.edu

8 Acknowledgments

The authors like to thank the NML working group members for their patience.

9 Intellectual Property Statement

The OGF takes no position regarding the validity or scope of any intellectual property or other rights that might be claimed to pertain to the implementation or use of the technology described in this document or the extent to which any license under such rights might or might not be available; neither does it represent that it has made any effort to identify any such rights. Copies of claims of rights made available for publication and any assurances of licenses to be made available, or the result of an attempt made to obtain a general license or permission for the use of such proprietary rights by implementers or users of this specification can be obtained from the OGF Secretariat.

The OGF invites any interested party to bring to its attention any copyrights, patents or patent applications, or other proprietary rights which may cover technology that may be required to practice this recommendation. Please address the information to the OGF Executive Director.

10 Disclaimer

This document and the information contained herein is provided on an “As Is” basis and the OGF disclaims all warranties, express or implied, including but not limited to any warranty that the use of the information herein will not infringe any rights or any implied warranties of merchantability or fitness for a particular purpose.

11 Full Copyright Notice

Copyright © Open Grid Forum (2008-2012). Some Rights Reserved.

This document and translations of it may be copied and furnished to others, and derivative works that comment on or otherwise explain it or assist in its implementation may be prepared, copied, published and distributed, in whole or in part, without restriction of any kind, provided that the above copyright notice and this paragraph are included on all such copies and derivative works. However, this document itself may not be modified in any way, such as by removing the copyright notice or references to the OGF or other organizations, except as needed for the purpose of developing Grid Recommendations in which case the procedures for copyrights defined in the OGF Document process must be followed, or as required to translate it into languages other than English.

The limited permissions granted above are perpetual and will not be revoked by the OGF or its successors or assignees.

References

Normative References

[URN-OGF-NETWORK] Freek Dijkstra. URN:OGF:network specification. GWD-R-P (Work in Progress), May 2011.

[RFC 2119] Scott Bradner. Key words for use in RFCs to Indicate Requirement Levels. RFC 2119 (Best Current Practice), March 1997. URL <http://tools.ietf.org/html/rfc2119>.

[RFC 3986] Tim Berners-Lee, Roy T. Fielding, and Larry Masinter. Uniform Resource Identifier (URI): Generic Syntax RFC 3986 (Standards Track), January 2005. URL <http://tools.ietf.org/html/rfc3986>.

Informative References