



# Grid Laboratory Of Wisconsin (GLOW)

Sridhara Dasu, Dan Bradley, Steve Rader  
Department of Physics

Miron Livny, Sean Murphy, Erik Paulson  
Department of Computer Science



<http://www.cs.wisc.edu/condor/glow>

# UW “Enterprise Level” Grid

- Condor pools at various departments integrated into a campus wide grid
  - **Grid Laboratory of Wisconsin**
- Older private Condor pools at other departments
  - ~1000 ~1GHz Intel CPUs at CS
  - ~100 ~2GHz Intel CPUs at Physics
  - ...
- **Condor jobs flock from various departments to GLOW**
- **Excellent utilization**
  - Especially when the Condor Standard Universe is used
    - Preemption, Checkpointing, Job Migration



# Grid Laboratory of Wisconsin

2003 Initiative funded by NSF/UW

Six GLOW Sites

- Computational Genomics, Chemistry
- Amanda, Ice-cube, Physics/Space Science
- High Energy Physics/CMS, Physics
- Materials by Design, Chemical Engineering
- Radiation Therapy, Medical Physics
- Computer Science


GLOW phases-1,2 + non-GLOW funded nodes  
already have ~1000 Xeons + 100 TB disk

# Condor/GLOW Ideas

- Exploit commodity hardware for high throughput computing
  - The base hardware is the same at all sites
  - Local configuration optimization as needed
    - e.g., Number of CPU elements vs storage elements
  - Must meet global requirements
    - It turns out that our initial assessment calls for almost identical configuration at all sites
- Managed locally at 6 sites
  - Shared globally across all sites
  - Higher priority for local jobs
- Neighborhood association style
  - Cooperative planning, operations ...


# GLOW Deployment

- **GLOW Phase-I and II are Commissioned**
  - CPU
    - 66 nodes each @ ChemE, CS, LMCG, MedPhys, Physics
    - 30 nodes @ IceCube
    - ~100 extra nodes @ CS (50 ATLAS + 50 CS)
    - 26 extra nodes @ Physics
    - **Total CPU: ~1000**
  - Storage
    - Head nodes @ at all sites
    - 45 TB each @ CS and Physics
    - **Total storage: ~ 100 TB**
- **GLOW Resources are used at 100% level**
  - Key is to have multiple user groups
- **GLOW continues to grow**



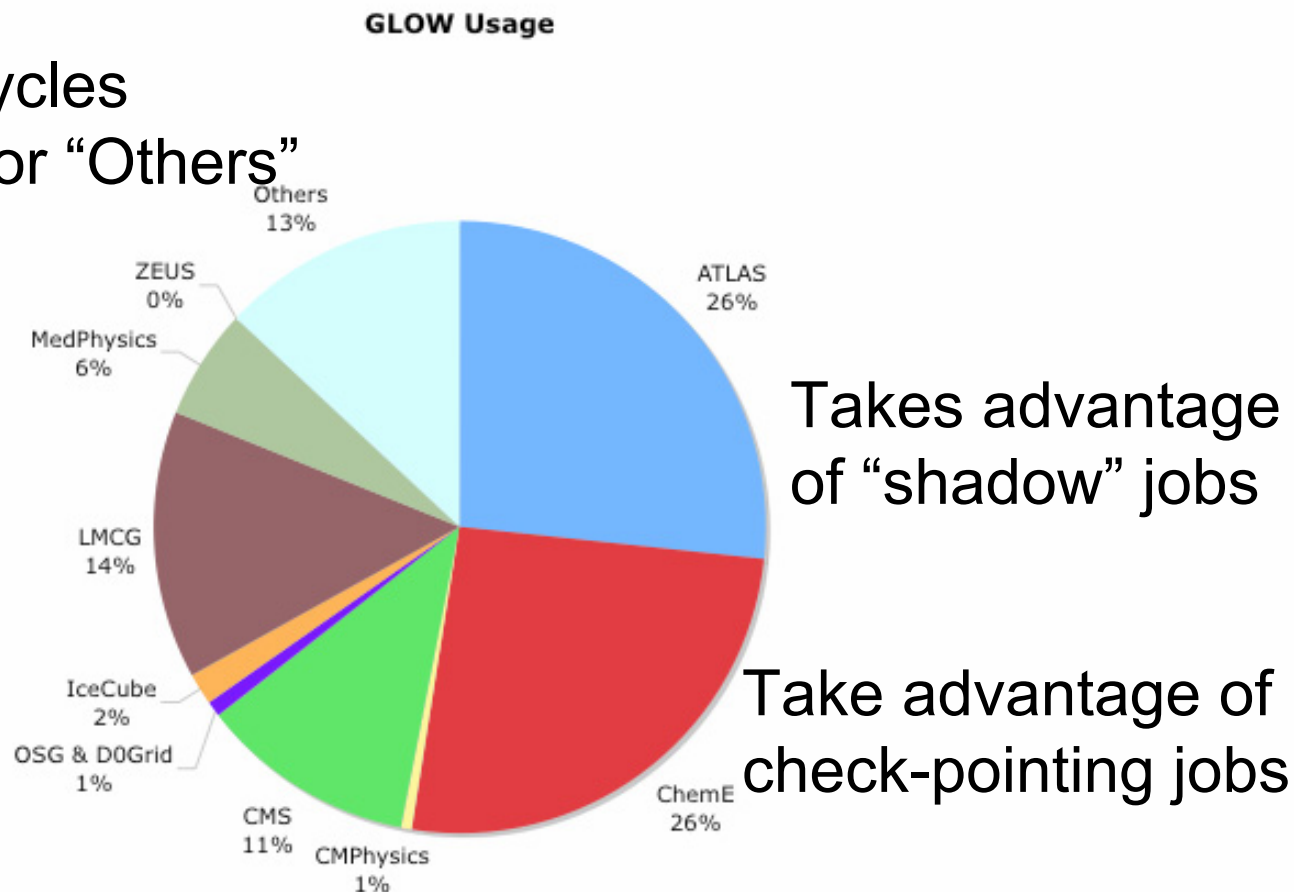
# GLOW Usage

- GLOW Nodes are always running hot!
  - CS + Guests
    - Largest user
    - Serving guests - many cycles delivered to guests!
  - ChemE
    - Largest community
  - HEP/CMS
    - Production for collaboration
    - Production and analysis of local physicists
  - LMCG
    - Standard Universe
  - Medical Physics
    - MPI jobs
  - IceCube
    - Simulations



# GLOW Usage Since February 2004

Leftover cycles  
available for “Others”



Over 7.6 million CPU-Hours (865 CPU-Years) served!

# Example Uses


- ATLAS
  - Over 15 Million proton collision events simulated at 10 minutes each
- CMS
  - Over 10 Million events simulated in a month - many more events reconstructed and analyzed
- Computational Genomics
  - Prof. Schwartz asserts that GLOW has opened up new paradigm of work patterns in his group
    - They no longer think about how long a particular computational job will take - they just do it
- Chemical Engineering
  - Students do not know where the computing cycles are coming from - they just do it



# CMS Usage of UW Campus Grid


Seamless support  
for bursts of activity:  
Opportunistic use of  
friendly resources

Data intensive jobs  
Local CMS-GLOW  
storage w/ campus-  
wide computing  
elements - high speed  
campus network



# ATLAS Use of GLOW

- UW ATLAS group is sold on GLOW
  - First new member of GLOW



- Efficiently used idle resources
- Used suspension mechanism to keep jobs in background when higher priority “owner” jobs kick-in

# New GLOW Members

- Proposed minimum involvement
  - One rack with about 50 CPUs
- Identified system support person who joins GLOW-tech
  - Can be an existing member of GLOW-tech
- PI joins the GLOW-exec
- Adhere to current GLOW policies
- Sponsored by existing GLOW members
  - UW ATLAS group and other physics groups were proposed by CMS and CS, and were accepted as new members
    - UW ATLAS using bulk of GLOW cycles (housed @ CS)
  - Expressions of interest from other groups

# GLOW & Condor Development

GLOW presents distributed computing researchers with an ideal laboratory of real users with diverse requirements

- Early commissioning and stress testing of new Condor releases in an environment controlled by Condor team
  - Results in robust releases for world-wide deployment
- New features in Condor Middleware, examples:
  - Group wise or hierarchical priority setting
  - Rapid-response with large resources for short periods of time for high priority interrupts
  - Hibernating shadow jobs instead of total preemption (HEP cannot use Standard Universe jobs)
  - MPI use (Medical Physics)
  - Condor-C (High Energy Physics and Open Science Grid)

# Open Science Grid & GLOW

- OSG Jobs can run on GLOW
  - Gatekeeper routes jobs to local condor cluster
  - Jobs flock to campus wide, including the GLOW resources
  - dCache storage pool is also a registered OSG storage resource
  - Beginning to see some use
- Now actively working on rerouting GLOW jobs to the rest of OSG
  - Users do NOT have to adapt to OSG interface and separately manage their OSG jobs
  - New Condor code development

# Summary

- Wisconsin campus grid, GLOW, has become an indispensable computational resource for several domain sciences
- Cooperative planning of acquisitions, installation and operations results in large savings
- Domain science groups no longer worry about setting up computing - they do their science!
  - Empowers individual scientists
  - Therefore, GLOW is growing on our campus
- By pooling together our resources we are able to harness larger than our *individual-share* at times of critical need to produce science results in a timely way
- Provides a working laboratory for computer science studies